Model Checking classical mutual exclusion algorithms

Ivan Ivchenko

01/08/2024

0.1   Brief introduction

During this project we will take some of the mutual exclusion algorithms and perform a model-checking for them. The algorithms that were selected:

1. Dekker's algorithm

2. Peterson's algorithm

The tool that we are going to use is UPPAAL. The project and all of the files will also be published on my GitHub account (link).

0.2   Dekker's Algorithm

Dekker's algorithm is the first known correct solution to the mutual exclusion problem in concurrent programming where processes only communicate via shared memory.

 If two processes try to enter a critical section at the same time, the algorithm will only allow one of them to do so, based on whose turn it is at that moment. If one process has already entered the critical section, the other will wait for the first to leave. This is implemented by using two flags (indicators of "intent" to enter the critical section) and a turn variable (indicating which process has queued).

 We can describe this using pseudocode:

Algorithm 1: Initialization

---
1 flag[0] := false
2 flag[1] := false
3 turn := 0 // or 1

Algorithm 2: Algorithm for process A	**Algorithm 3:** Algorithm for process B
1 flag[0] := true	1 flag[1] := true
2 **while** *flag[1] = true* **do**	2 **while** *flag[0] = true* **do**
3 **if** *turn = 1* **then**	3 **if** *turn = 0* **then**
4 flag[0] := false	4 flag[1] := false
5 **while** *turn = 1* **do**	5 **while** *turn = 0* **do**
6 busy-wait	6 busy-wait
7 **end**	7 **end**
8 flag[0] := true	8 flag[1] := true
9 **end**	9 **end**
10 **end**	10 **end**
11 // critical section	11 // critical section
12 ...	12 ...
13 turn := 1	13 turn := 0
14 flag[0] := false	14 flag[1] := false
15 // end of critical section	15 // end of critical section
16 ...	16 ...

Dekker's algorithm ensures mutual exclusion and prevents deadlock. If process A remains in the "while flag[1]" loop indefinitely, eventually process B will enter its critical section and set turn = 0. A will then exit its inner "while turn = 1" loop, set flag[0] = true, and wait for flag[1] to become false. When B tries to re-enter the critical section, it will set flag[1] = false and loop on "while turn = 0." This allows A to proceed into its critical section. Removing the "turn = 0" check in A's loop would introduce the possibility of starvation, so all steps are crucial.

Because we are using UPPAAL GUI, in the next images I will present the Dekker's algorithm loaded into the tool and how i defined it.

Before we begin we need to declare global variables and our system specifications.

```
1  bool flag[2] =
       {false, false};
2  int turn = 1;
```

Figure 1: Global variables declarations

```
1  P0 = proc(0);
2  P1 = proc(1);
3  system P0, P1;
```

Figure 2: System declarations

```
1  const int pid
```

Figure 3: Process parameter

Now moving on to the process definition (Figure 4). For this, our model with states (locations) that correspond to actions in the pseudocode, the guard statements (what should hold in order to move from one state to another) are marked green and the update statements (what values will be changed after the step) are marked blue. For better understanding (Algorithm 4):
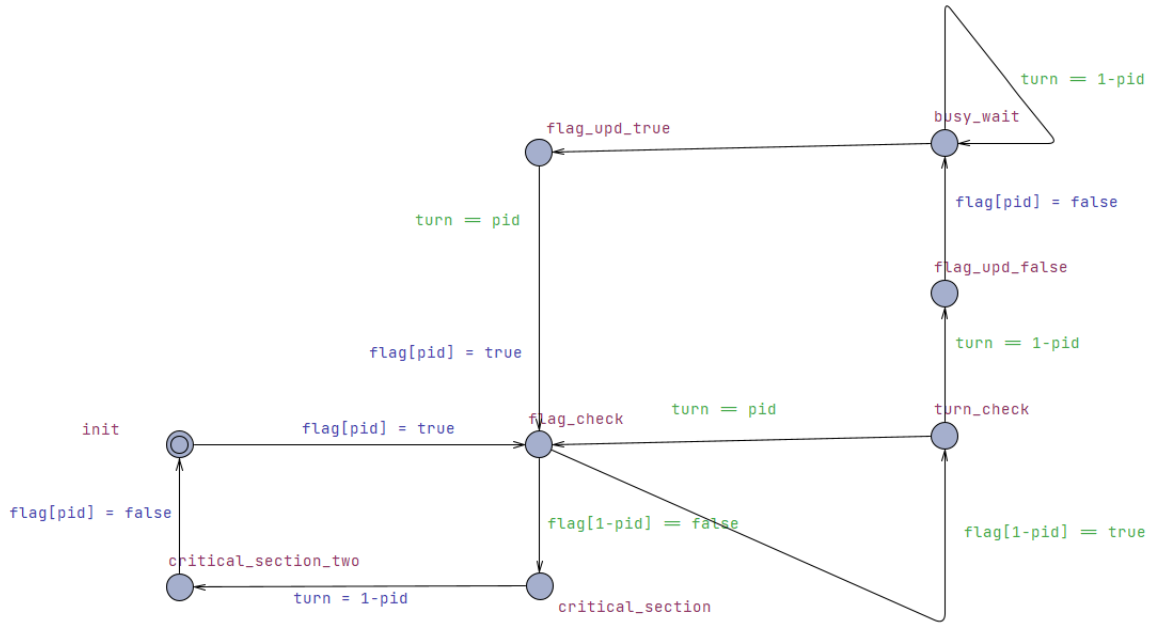

Figure 4: Definition of our algorithm

Algorithm 4: Algorithm interpretation

1 (init)
2 flag[1] := true
3 (flag_check)
4 while *flag[0] = true* do
5   (turn_check)
6   if *turn = 0* then
7   (flag_upd_false)
8   flag[1] := false
9   (busy_wait)
10   while *turn = 0* do
11   busy-wait
12   end
13   (flag_upd_true)
14   flag[1] := true
15   end
16 end
17 //critical section
18 (critical_section)
19 turn := 0
20 (critical_section_two)
21 flag[1] := false
22 // end of critical section
23 (init)

For the properties to be verified we specify:

1. $A[]$ not ($P0$.critical_section $\land$ $P1$.critical_section) $\land$
   not ($P0$.critical_section_two $\land$ $P1$.critical_section) $\land$
   not ($P1$.critical_section_two $\land$ $P0$.critical_section) $\land$
   not ($P1$.critical_section_two $\land$ $P0$.critical_section_two) - with UPPAAL specification we use $A[]$ symbol to represent "on all paths at all points in time" or we can rephrase it to "always", and because with our specification of the critical section that consists of 2 states (locations), we need to check for all possible combination of states that will trigger mutual exclusion

2. $A[]$ not deadlock - with UPPAAL specification we use A[] symbol to represent "on all paths at all points in time" or we can rephrase it to "always", this ensures that our model is defined and initialized correctly

   After this, all we have to do is to run a check to see if our properties are being satisfied.

Status

UPPAAL version 5.0.0 (rev. 714BA9DB36F49691), June 2023 -- server. Licensed
A[] not (P0.critical_section and P1.critical_section) && not (P0.critical_section_
Verification/kernel/elapsed time used: 0s / 0.016s / 0.002s.
Resident/virtual memory usage peaks: 17,784KB / 63,092KB.
Property is satisfied.
A[] not deadlock
Verification/kernel/elapsed time used: 0s / 0s / 0.004s.
Resident/virtual memory usage peaks: 17,916KB / 63,312KB.
Property is satisfied.

Figure 5: Results of our check

0.3 Peterson's Algorithm (for 2 processes)

Peterson's algorithm (or Peterson's solution) is a concurrent programming algorithm for mutual exclusion that allows two or more processes to share a single-use resource without conflict, using only shared memory for communication.

The algorithm uses two variables: flag and turn. A flag[n] value of true indicates that the process n wants to enter the critical section. Entrance to the critical section is granted for process A if B does not want to enter its critical section or if B has given priority to A by setting turn to 0.

Using the pseudocode , we can present Peterson's algorithm like this:

Algorithm 5: Initialization
---
1 bool flag[2] = false, false;
2 int turn;

Algorithm 6: Algorithm for process A	**Algorithm 7:** Algorithm for process B
1 flag[0] = true;	**1** flag[1] = true;
2 turn = 1;	**2** turn = 0;
3 **while** *flag[1] == true and turn==1* **do**	**3** **while** *flag[0] == true and turn==0* **do**
4 $\quad\mid\quad$ busy-wait	**4** $\quad\mid\quad$ busy-wait
5 **end**	**5** **end**
6 //critical section	**6** //critical section
7	**7** .....
8 //end of critical section	**8** //end of critical section
9 flag[0] = false;	**9** flag[1] = false;

A and B can never be in the critical section at the same time. If A is in its critical section, then flag[0] is true. In addition, either flag[1] is false (meaning that B has left its critical section), or turn is 0 (meaning that B is just now trying to enter the critical section, but graciously waiting). So if both processes are in their critical sections, then we conclude that the state must satisfy flag[0] and flag[1] and turn = 0 and turn = 1. No state can satisfy both turn = 0 and turn = 1, so there can be no state where both processes are in their critical sections.

Because we are using UPPAAL GUI, in the next images I will present the Peterson's algorithm loaded into the tool and how i defined it.

```
1  bool flag[2] =
       {false, false};
2
3  int[0,1] turn;
```

```
1  P0 = proc(0);
2  P1 = proc(1);
3  system P0, P1;
```

```
1  const int pid
```

Figure 8: Process parameter

Figure 6: Global variables declarations

Figure 7: System declarations

Now moving on to the process definition. For this, our model with states (locations) that correspond to actions in the pseudocode. For better understanding:

Algorithm 8: Algorithm for process A

---

1 (init)
2 flag[0] = true;
3 (turn_upd)
4 turn = true;
5 (flag_and_turn_check)
6 while *flag[1] == true* **and** *turn == 1* **do**
7 │ **(busy_wait)**
8 end
9 // **critical section**
10 (critical_section)
11 // **end of critical section**
12 flag[0] = false;
13 (init)

Now moving on to the process definition. For this I've simplified our model for a better understanding of the concept of algorithm, the guard statements (what should hold in order to move from one state to another) are marked green and the update statements (what values will be changed after the step) are marked blue (Figure 9).

Figure 9: Definition of our algorithm

For the properties to be verified we specify:

1. $A[]$ not ($P0$.critical_section $\wedge$ $P1$.critical_section) - with UPPAAL specification we use A[] symbol to represent "on all paths at all points in time" or we can rephrase it to "always", with this we check for the mutual exclusion.

2. *A[]* not deadlock - with UPPAAL specification we use A[] symbol to represent "on all paths at all points in time" or we can rephrase it to "always", this ensures that our model is defined and initialized correctly

   After this, all we have to do is to run a check to see if our properties are being satisfied.

images/Peterson_check.png

Figure 10: Results of our check

0.4 Conclusion

In this project, we successfully modeled and verified Dekker's and Peterson's mutual exclusion algorithms using UPPAAL. Through the creation of process templates, the definition of shared variables, and the formulation of transitions and guard conditions, we were able to represent the core logic of these algorithms accurately.

We verified critical properties such as mutual exclusion and absence of deadlock by utilizing UPPAAL's query capabilities. The mutual exclusion property ensured that no two processes could be in the critical section simultaneously, while the absence of deadlock guaranteed that processes would not get stuck indefinitely waiting to enter the critical section. Through this project, we learned how to use the tool for model-checking and also familiarized ourselves with the basic mutual exclusion algorithms.